

2016 Election Prediction

Jorge Sifuentes & Pedro Aristizabal (PSTAT 131)

5/21/2019

Predicting voter behavior is complicated for many reasons despite the tremendous effort in collecting, analyzing, and understanding many available datasets. For our final project, we will analyze the 2016 presidential election dataset, but, first, some background.

Background

The presidential election in 2012 did not come as a surprise. Some correctly predicted the outcome of the election correctly including Nate Silver, and many speculated his approach.

Despite the success in 2012, the 2016 presidential election came as a big surprise to many, and it was a clear example that even the current state-of-the-art technology can surprise us.

Answer the following questions in one paragraph for each.

1. What makes voter behavior prediction (and thus election forecasting) a hard problem?

For one there are multiple dimensions to consider (location, age, education level, etc ...). There is also the number of candidates to consider and also how candidate popularity is sensitive to unforeseen real time change (A candidate's scandal gets uncovered for instance). It's not obvious how some factors may correlate. The statistics may be compromised as politicians in that position tend to be nefarious in nature. Electoral College is complicated.

2. What was unique to Nate Silver's approach in 2012 that allowed him to achieve good predictions?

Nate Silver prediction model was unique in the following ways: FiveThirtyEight's probabilities are based on the accuracy of polling averages. In addition to a systematic national polling error, they also simulate potential errors across regional or demographic lines. They assume that polling errors are correlated. Their historical models and the recent election show that making inferences from early voting is generally not good.

3. What went wrong in 2016? What do you think should be done to make future predictions better?

From the article the errors in predictions were coming from wrong assumptions and over generalizations, such as state vs national polls. It seems that more information can be accurately derived based on historical trends of polls and election outcomes. The accuracy seems to not come directly from the model but how well we can predict the error. Some challenges for future elections is predicting trends with late undecided voters, and the changing of party preferences of some states such as Arizona and Pennsylvania. Journalist.

Data

```
election.raw = read.csv("election.csv") %>% as.tbl
census_meta = read.csv("metadata.csv", sep = ";") %>% as.tbl
census = read.csv("census.csv") %>% as.tbl
census$CensusTract = as.factor(census$CensusTract)
```

Election data

Following is the first few rows of the `election.raw` data:

| county | fips | candidate | state | votes |
|--------|------|-----------------|-------|----------|
| NA | US | Donald Trump | US | 62984825 |
| NA | US | Hillary Clinton | US | 65853516 |
| NA | US | Gary Johnson | US | 4489221 |
| NA | US | Jill Stein | US | 1429596 |
| NA | US | Evan McMullin | US | 510002 |
| NA | US | Darrell Castle | US | 186545 |

The meaning of each column in `election.raw` is clear except `fips`. The acronym is short for Federal Information Processing Standard.

In our dataset, `fips` values denote the area (US, state, or county) that each row of data represent: i.e., some rows in `election.raw` are summary rows. These rows have `county` value of `NA`. There are two kinds of summary rows:

- Federal-level summary rows have `fips` value of `US`.
- State-level summary rows have names of each states as `fips` value.

Census data

Following is the first few rows of the `census` data:

| CensusTract | State | County | TotalPop | Men | Women | Hispanic | White | Black | Native | Asian | Pacific | C |
|-------------|---------|---------|----------|------|-------|----------|-------|-------|--------|-------|---------|---|
| 1001020100 | Alabama | Autauga | 1948 | 940 | 1008 | 0.9 | 87.4 | 7.7 | 0.3 | 0.6 | 0.0 | |
| 1001020200 | Alabama | Autauga | 2156 | 1059 | 1097 | 0.8 | 40.4 | 53.3 | 0.0 | 2.3 | 0.0 | |
| 1001020300 | Alabama | Autauga | 2968 | 1364 | 1604 | 0.0 | 74.5 | 18.6 | 0.5 | 1.4 | 0.3 | |
| 1001020400 | Alabama | Autauga | 4423 | 2172 | 2251 | 10.5 | 82.8 | 3.7 | 1.6 | 0.0 | 0.0 | |
| 1001020500 | Alabama | Autauga | 10763 | 4922 | 5841 | 0.7 | 68.5 | 24.8 | 0.0 | 3.8 | 0.0 | |
| 1001020600 | Alabama | Autauga | 3851 | 1787 | 2064 | 13.1 | 72.9 | 11.9 | 0.0 | 0.0 | 0.0 | |

Census data: column metadata

Column information is given in `metadata`.

| CensusTract | Census.tract.ID | numeric |
|-------------|---|---------|
| State | State, DC, or Puerto Rico | string |
| County | County or county equivalent | string |
| TotalPop | Total population | numeric |
| Men | Number of men | numeric |
| Women | Number of women | numeric |
| Hispanic | % of population that is Hispanic/Latino | numeric |
| White | % of population that is white | numeric |
| Black | % of population that is black | numeric |
| Native | % of population that is Native American or Native Alaskan | numeric |
| Asian | % of population that is Asian | numeric |
| Pacific | % of population that is Native Hawaiian or Pacific Islander | numeric |
| Citizen | Number of citizens | numeric |
| Income | Median household income (\$) | numeric |
| IncomeErr | Median household income error (\$) | numeric |

| CensusTract | Census.tract.ID | numeric |
|-----------------|---|---------|
| IncomePerCap | Income per capita (\$) | numeric |
| IncomePerCapErr | Income per capita error (\$) | numeric |
| Poverty | % under poverty level | numeric |
| ChildPoverty | % of children under poverty level | numeric |
| Professional | % employed in management, business, science, and arts | numeric |
| Service | % employed in service jobs | numeric |
| Office | % employed in sales and office jobs | numeric |
| Construction | % employed in natural resources, construction, and maintenance | numeric |
| Production | % employed in production, transportation, and material movement | numeric |
| Drive | % commuting alone in a car, van, or truck | numeric |
| Carpool | % carpooling in a car, van, or truck | numeric |
| Transit | % commuting on public transportation | numeric |
| Walk | % walking to work | numeric |
| OtherTransp | % commuting via other means | numeric |
| WorkAtHome | % working at home | numeric |
| MeanCommute | Mean commute time (minutes) | numeric |
| Employed | % employed (16+) | numeric |
| PrivateWork | % employed in private industry | numeric |
| PublicWork | % employed in public jobs | numeric |
| SelfEmployed | % self-employed | numeric |
| FamilyWork | % in unpaid family work | numeric |
| Unemployment | % unemployed | numeric |

Data wrangling

4. Remove summary rows from `election.raw` data: i.e.,

- Federal-level summary into a `election_federal`.
- State-level summary into a `election_state`.
- Only county-level data is to be in `election`.

```
election_federal <- election.raw %>% filter(fips == "US")
election_state <- election.raw %>% filter(fips != "US" & is.na(county) & fips != "2000" & fips != "4610")
election <- election.raw %>% filter(!is.na(county))

##Federal level summary
kable(election_federal %>% head)
```

| county | fips | candidate | state | votes |
|--------|------|-----------------|-------|----------|
| NA | US | Donald Trump | US | 62984825 |
| NA | US | Hillary Clinton | US | 65853516 |
| NA | US | Gary Johnson | US | 4489221 |
| NA | US | Jill Stein | US | 1429596 |
| NA | US | Evan McMullin | US | 510002 |
| NA | US | Darrell Castle | US | 186545 |

```
##State level summary
kable(election_state%>% head)
```


| county | fips | candidate | state | votes |
|--------|------|-----------------|-------|---------|
| NA | CA | Hillary Clinton | CA | 8753788 |
| NA | CA | Donald Trump | CA | 4483810 |
| NA | CA | Gary Johnson | CA | 478500 |
| NA | CA | Jill Stein | CA | 278657 |
| NA | CA | Gloria La Riva | CA | 66101 |
| NA | FL | Donald Trump | FL | 4617886 |

```
#county level data only
kable(election %>% head)
```

| county | fips | candidate | state | votes |
|--------------------|-------|-----------------|-------|---------|
| Los Angeles County | 6037 | Hillary Clinton | CA | 2464364 |
| Los Angeles County | 6037 | Donald Trump | CA | 769743 |
| Los Angeles County | 6037 | Gary Johnson | CA | 88968 |
| Los Angeles County | 6037 | Jill Stein | CA | 76465 |
| Los Angeles County | 6037 | Gloria La Riva | CA | 21993 |
| Cook County | 17031 | Hillary Clinton | IL | 1611946 |

5. How many named presidential candidates were there in the 2016 election? Draw a bar chart of all votes received by each candidate

```
ggplot(election_federal, aes(x = candidate, y = votes,)) + geom_bar(stat="identity")+coord_flip()
```



As We can see, there were 32 candidates for the 2016 presidential elections.

6. Create variables `county_winner` and `state_winner` by taking the candidate with the highest proportion of votes. Hint: to create `county_winner`, start with `election`, group by `fips`, compute `total` votes, and `pct = votes/total`. Then choose the highest row using `top_n` (variable `state_winner` is similar).

```
election <- election %>% group_by(fips) %>% mutate(total = sum(votes))
election <- election %>% group_by(fips) %>% mutate(pct = votes/total)

## County winner variable
county_winner <- election %>% group_by(fips) %>% top_n(1)

## Selecting by pct
election_state <- election_state %>% group_by(fips) %>% mutate(total = sum(votes))
election_state <- election_state %>% group_by(fips) %>% mutate(pct = votes/total)

#state winner variable
state_winner <- election_state %>% group_by(fips) %>% top_n(1)

## Selecting by pct
```

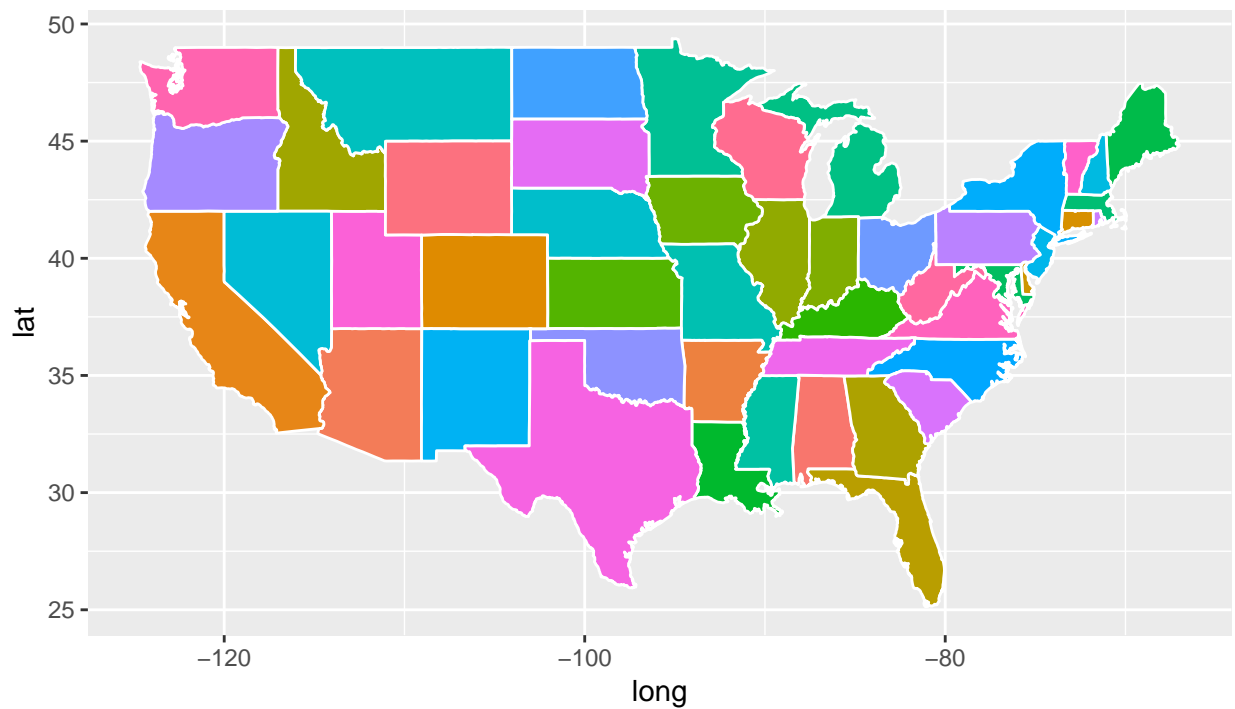
Visualization

Visualization is crucial for gaining insight and intuition during data mining. We will map our data onto maps.

The R package `ggplot2` can be used to draw maps. Consider the following code.

```
states = map_data("state")

ggplot(data = states) +
  geom_polygon(aes(x = long, y = lat, fill = region, group = group), color = "white") +
  coord_fixed(1.3) +
  guides(fill=FALSE) # color legend is unnecessary and takes too long
```

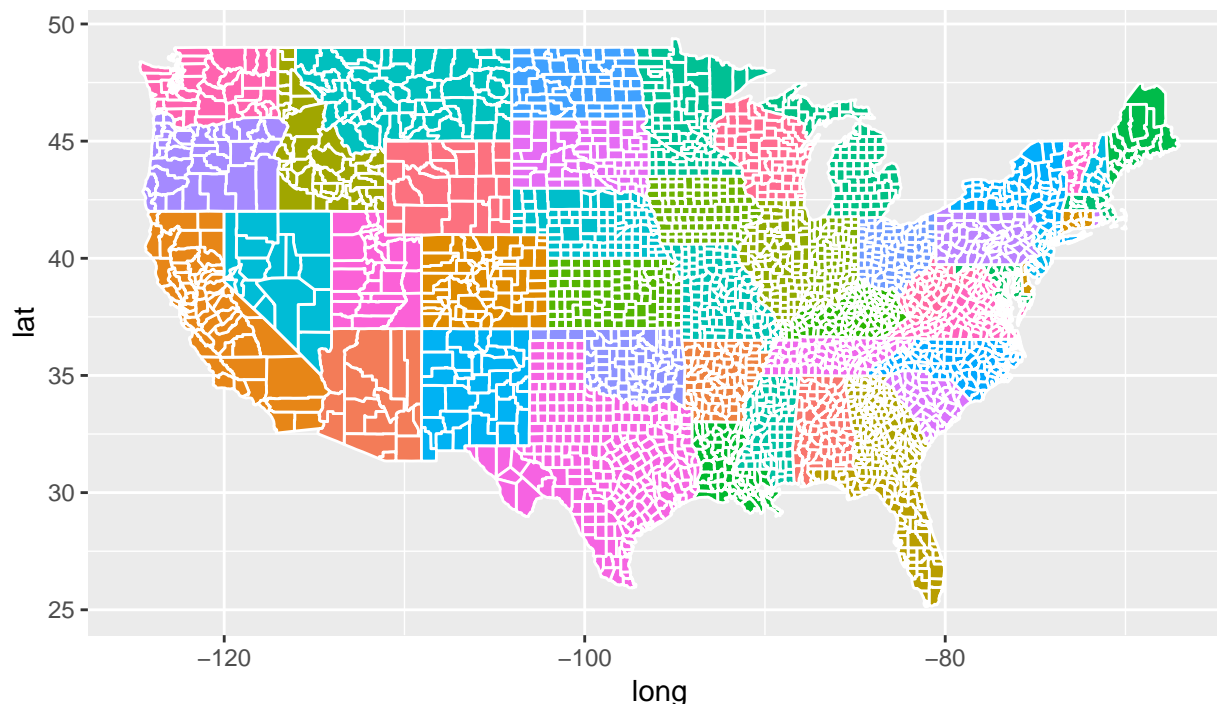


The variable `states` contain information to draw white polygons, and fill-colors are determined by `region`.

7. Draw county-level map by creating `counties = map_data("county")`. Color by county

```
counties = map_data("county")

ggplot(data = counties) +
  geom_polygon(aes(x = long, y = lat, fill = region, group = group), color = "white") +
  coord_fixed(1.3) +
  guides(fill=FALSE)
```



8. Now color the map by the winning candidate for each state. First, combine `states` variable and `state_winner` we created earlier using `left_join()`. Note that `left_join()` needs to match up values of states to join the tables; however, they are in different formats: e.g. `AZ` vs. `arizona`. Before using `left_join()`, create a common column by creating a new column for `states` named `fips` = `state.abb[match(some_column, some_function(state.name))]`. Replace `some_column` and `some_function` to complete creation of this new column. Then `left_join()`.

```
fips = state.abb[match(states$region, tolower(state.name))]
```

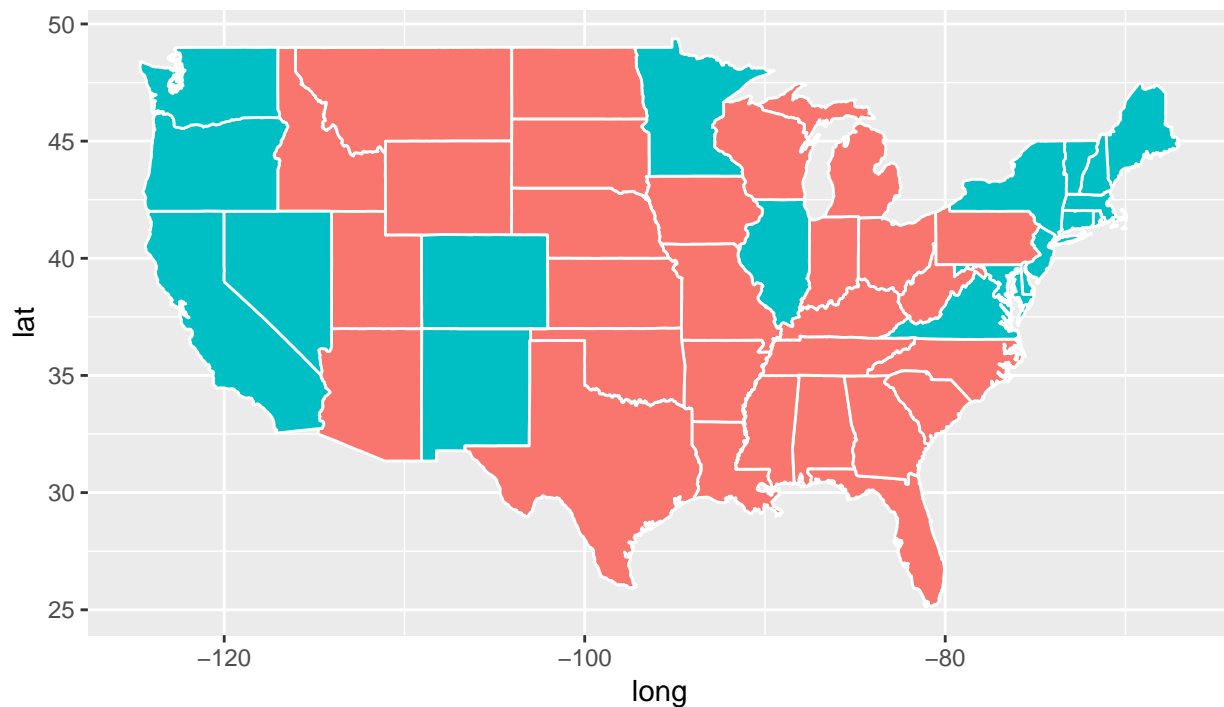
```
states <- cbind(states, fips)
```

```
states <- left_join(state_winner, states)
```

```
## Joining, by = "fips"
```


```
## Warning: Column `fips` joining factors with different levels, coercing to
## character vector
```

```
ggplot(data = states) +
  geom_polygon(aes(x = long, y = lat, fill = candidate, group = group), color = "white") +
  coord_fixed(1.3) +
  guides(fill=FALSE)
```



9. The variable county does not have fips column. So we will create one by pooling information from maps::county.fips. Split the polynome column to region and subregion. Use left_join() combine county.fips into county. Also, left_join() previously created variable county_winner.

```
counties <- left_join( maps::county.fips %>% separate(polynome, c("region","subregion"), sep = ",", ext
## Joining, by = c("region", "subregion")
counties$fips <- as.factor(counties$fips)
counties <- left_join(county_winner,counties)
## Joining, by = "fips"
## Warning: Column `fips` joining factors with different levels, coercing to
## character vector
ggplot(data = counties) +
  geom_polygon(aes(x = long, y = lat, fill = candidate, group = group), color = "white") +
  coord_fixed(1.3) +
  guides(fill=FALSE)
```


10. Create a visualization of your choice using census data. Many exit polls noted that demographics played a big role in the election. Use this Washington Post article and this R graph gallery for ideas and inspiration.

For this question, we decided to create a barplot showing the amount of non-citizens found in each state, looking to find which states have the highest population. This info might be helpful to notice voting patterns for future elections.

```
census <- census %>% group_by(County) %>% mutate(Noncitizens = (TotalPop - Citizen))

census <- census %>% group_by(County) %>% mutate(Noncitizen_percentage = Noncitizens/TotalPop)

ggplot(census, aes(x = State, y = Noncitizens)) + geom_bar(stat="identity") + coord_flip()
```



11. The census data contains high resolution information (more fine-grained than county-level). In this problem, we aggregate the information into county-level data by computing TotalPop-weighted average of each attributes for each county. Create the following variables:

```
* _Clean census data `census.del`_:
start with `census`, filter out any rows with missing values,
convert {`Men`, `Employed`, `Citizen`} attributes to a percentages (meta data seems to be inaccurate)
compute `Minority` attribute by combining {Hispanic, Black, Native, Asian, Pacific}, remove {`Walk`,
_Many columns seem to be related, and, if a set that adds up to 100%, one column will be deleted._

* _Sub-county census data, `census.subct`_:
start with `census.del` from above, `group_by()` two attributes {`State`, `County`},
use `add_tally()` to compute `CountyTotal`. Also, compute the weight by `TotalPop/CountyTotal`.

* _County census data, `census.ct`_:
start with `census.subct`, use `summarize_at()` to compute weighted sum

* _Print few rows of `census.ct`_:
census = read.csv("census.csv") %>% as.tbl
census$CensusTract = as.factor(census$CensusTract)

census.del <- census %>% filter(complete.cases(census))

census.del <- census.del %>% mutate(Men = Men/TotalPop)
```

```

census.del <- census.del %>% mutate(Employed = Employed/TotalPop)
census.del <- census.del %>% mutate(Citizen = Citizen/TotalPop)

census.del <- census.del %>% mutate(Minority = Hispanic + Black + Native + Asian + Pacific)

census.del <- census.del %>% select(-c(Walk, PublicWork, Construction))

census.del<- census.del %>% select(-Hispanic,-Black,-Native,-Asian,-Pacific,-Women)

census.subct <- census.del %>% group_by(State, County) %>% add_tally(TotalPop)

census.subct <- census.subct %>% mutate(CountyTotal = n)

census.subct<- census.subct %>% select(-n)

census.subct <- census.subct %>% mutate(weight = TotalPop/CountyTotal)

census.ct<-census.subct %>% summarise_at(vars(Men:CountyTotal), funs(weighted.mean(., weight)))

census.ct<-as.data.frame(census.ct)

census.subct<-as.data.frame(census.subct)

kable(census.ct) %>% head

```

| ## [1] "State | County | Men | White | Citizen |
|-----------------|---------|-----------|------------|-----------|
| ## [2] "----- | ----- | ----- | ----- | ----- |
| ## [3] "Alabama | Autauga | 0.4843266 | 75.7882273 | 0.7374912 |
| ## [4] "Alabama | Baldwin | 0.4884866 | 83.1026163 | 0.7569406 |
| ## [5] "Alabama | Barbour | 0.5382816 | 46.2315944 | 0.7691222 |
| ## [6] "Alabama | Bibb | 0.5341090 | 74.4998894 | 0.7739781 |

Dimensionality reduction

- Run PCA for both county & sub-county level data. Save the first two principle components PC1 and PC2 into a two-column data frame, call it `ct.pc` and `subct.pc`, respectively. What are the most prominent loadings?

```

#for Subcounty
pr.out_sub=prcomp(census.subct[,4:31],scale=TRUE)
subct.pc<-pr.out_sub$rotation[,1:2] ##gives out PC1 and PC2

max(abs(subct.pc)[1:28]) ##Value corresponds to IncomePerCap (PC1)

```

```
## [1] 0.3181199
```

```
max(abs(subct.pc)[29:56]) ## Value corresponds to Transit (PC2)
```

```
## [1] 0.3937091
```

```

#County level
pr.out_ct=prcomp(census.ct[,3:28],scale=TRUE)
ct.pc<-pr.out_ct$rotation[,1:2] ##gives out PC1 and PC2

max(abs(ct.pc)[1:26]) ## Value corresponds to IncomePerCar

```

```
## [1] 0.3530767
```

```
max(abs(ct.pc)[27:52]) ## Value corresponds to IncomeErr
```

```
## [1] 0.3145022
```

As we see, the most prominent loadings at the sub county level are IncomePerCap (PC1) and Transit (PC2). For the county level, we get IncomePerCar(PC1) and IncomeErr(PC2). # Clustering

13. With census.ct, perform hierarchical clustering using Euclidean distance metric complete linkage to find 10 clusters. Repeat clustering process with the first 5 principal components of ct.pc. Compare and contrast clusters containing San Mateo County. Can you hypothesize why this would be the case?

```
#standardizing census.ct
scaled_ct = scale(census.ct[, -c(1,2)], center=TRUE, scale=TRUE)
#evaluating distance w/ euclidean method
census.dist = dist(scaled_ct,method="euclidean")
```

```
set.seed(1)
census.hclust = hclust(census.dist) # complete linkage
census.hclust=cutree(census.hclust, 10) ## k=10
table(census.hclust)
```

```
## census.hclust
##      1      2      3      4      5      6      7      8      9     10
## 2632  501      6      7      5      1     11     13     38      4
```

```
#now for top 5 PC
```

```
scaled_ct_PCA<-scale(pr.out_ct$x[,1:5])
PCA_dist<-dist(scaled_ct_PCA)
census.hclust.PCA<- cutree(hclust(PCA_dist),k=10)
table(census.hclust.PCA)
```

```
## census.hclust.PCA
##      1      2      3      4      5      6      7      8      9     10
## 2441  525     97      6      8     31      5     18      7     80
```

Classification

In order to train classification models, we need to combine county_winner and census.ct data. This seemingly straightforward task is harder than it sounds. Following code makes necessary changes to merge them into election.cl for classification.

```
tmpwinner = county_winner %>% ungroup %>%
  mutate(state = state.name[match(state, state.abb)]) %>%          ## state abbreviations
  mutate_at(vars(state, county), tolower) %>%                    ## to all lowercase
  mutate(county = gsub(" county| columbia| city| parish", "", county)) ## remove suffixes
tmpcensus = census.ct %>% mutate_at(vars(State, County), tolower)

election.cl = tmpwinner %>%
  left_join(tmpcensus, by = c("state"="State", "county"="County")) %>%
  na.omit

## saves meta information to attributes
```

```
attr(election.cl, "location") = election.cl %>% select(c(county, fips, state, votes, pct))
election.cl = election.cl %>% select(-c(county, fips, state, votes, pct))
```

Using the following code, partition data into 80% training and 20% testing:

```
set.seed(10)
n = nrow(election.cl)
in.trn= sample.int(n, 0.8*n)
trn.cl = election.cl[ in.trn,]
tst.cl = election.cl[-in.trn,]
```

Using the following code, define 10 cross-validation folds:

```
set.seed(20)
nfold = 10
folds = sample(cut(1:nrow(trn.cl), breaks=nfold, labels=FALSE))
```

Using the following error rate function:

```
calc_error_rate = function(predicted.value, true.value){
  return(mean(true.value!=predicted.value))
}
records = matrix(NA, nrow=3, ncol=2)
colnames(records) = c("train.error", "test.error")
rownames(records) = c("tree", "knn", "lda")
```

Classification: native attributes

13. Decision tree: train a decision tree by `cv.tree()`. Prune tree to minimize misclassification. Be sure to use the folds from above for cross-validation. Visualize the trees before and after pruning. Save training and test errors to `records` variable.

```
tmpwinner = county_winner %>% ungroup %>%
  mutate(state = state.name[match(state, state.abb)]) %>%
  mutate_at(vars(state, county), tolower) %>%
  mutate(county = gsub(" county| columbia| city| parish", "", county))
tmpcensus = census.ct %>% mutate_at(vars(State, County), tolower)

election.cl = tmpwinner %>%
  left_join(tmpcensus, by = c("state"="State", "county"="County")) %>% na.omit

attr(election.cl, "location") = election.cl %>% select(c(county, fips, state, votes, pct))

election.cl = election.cl %>% select(-c(county, fips, state, votes, pct))

set.seed(10)
n = nrow(election.cl)
in.trn= sample.int(n, 0.8*n)
trn.cl = election.cl[ in.trn,]
tst.cl = election.cl[-in.trn,]

set.seed(20)
nfold = 10
folds = sample(cut(1:nrow(trn.cl), breaks=nfold, labels=FALSE))
```

```

calc_error_rate = function(predicted.value, true.value){ return(mean(true.value!=predicted.value))
}

records = matrix(NA, nrow=3, ncol=2)

colnames(records) = c("train.error", "test.error")
rownames(records) = c("tree", "knn", "lda")


#actual stuff

cv.tree <- tree(trn.cl$candidate~., trn.cl)
summary(cv.tree)

##
## Classification tree:
## tree(formula = trn.cl$candidate ~ ., data = trn.cl)
## Variables actually used in tree construction:
## [1] "Transit"      "White"        "Income"       "Unemployment"
## [5] "Production"   "total"        "Minority"     "OtherTransp"
## Number of terminal nodes: 12
## Residual mean deviance: 0.3612 = 882.8 / 2444
## Misclassification error rate: 0.06393 = 157 / 2456

##Un-pruned
library(rpart)
library("maptree")
draw.tree(cv.tree, nodeinfo=TRUE)

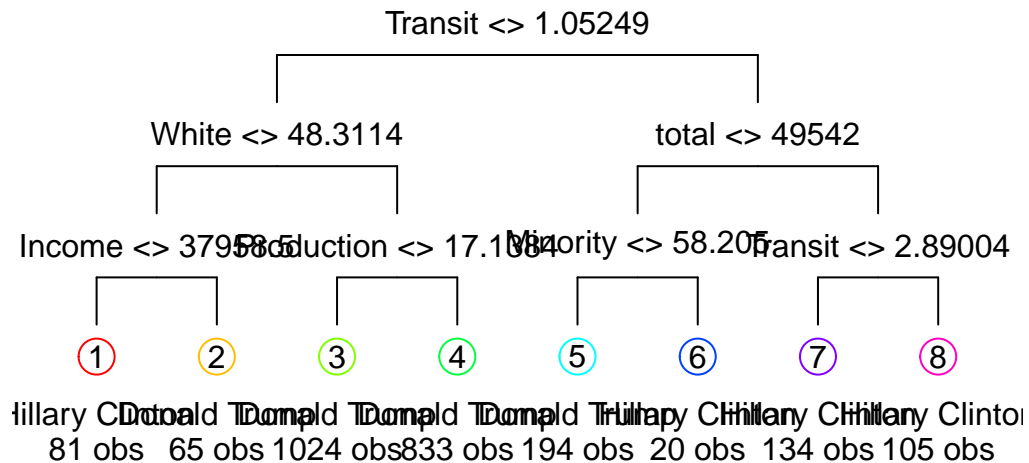
```



```
#pruned w K-cross means we have to use cv.tree()
cv = cv.tree(cv.tree, folds, FUN=prune.misclass, K=10)
best.cv = cv$size[which.min(cv$dev)]
best.cv
```

```
## [1] 8
```

```
pruned.tree<-prune.tree(cv.tree,best=best.cv)
draw.tree(pruned.tree)
```



```

prediction<- predict(pruned.tree,tst.cl,type='class')
test.error<-calc_error_rate(prediction, tst.cl$candidate)

train.prediction<-predict(pruned.tree,trn.cl, type='class')
train.error<-calc_error_rate(train.prediction,trn.cl$candidate)

records[1,]<-c(train.error,test.error)
records

```

```

##      train.error test.error
## tree  0.07247557 0.08469055
## knn           NA         NA
## lda           NA         NA

```

14. K-nearest neighbor: train a KNN model for classification. Use cross-validation to determine the best number of neighbors, and plot number of neighbors vs. resulting training and validation errors. Compute test error and save to records.

```

kvec = c(2, seq(10,100, length.out=10))

do.chunk <- function(chunkid, folddef, Xdat, Ydat, k){
  train = (folddef!=chunkid)
  Xtr = Xdat[train,]
  Ytr = Ydat[train]

  Xvl = Xdat[!train,]
  Yvl = Ydat[!train]

```



```

## get classifications for current training chunks
predYtr = knn(train = Xtr, test = Xtr, cl = Ytr, k = k)
## get classifications for current test chunk
predYvl = knn(train = Xtr, test = Xvl, cl = Ytr, k = k)

data.frame(train.error = calc_error_rate(predYtr, Ytr),
           val.error = calc_error_rate(predYvl, Yvl))
}

tmpDF <- data.frame(neighbors = 0, avg.train.error = 0,
                   avg.val.error = 0)

for (i in kvec){

  tmp = ldply(1:nfold, do.chunk, # Apply do.chunk() function to each fold
             folddef=folds, Xdat=trn.cl[3:28], Ydat=trn.cl$candidate , k=i)

  avg.val.error = mean(tmp$val.error)
  avg.train.error = mean(tmp$train.error)

  tmpDF <- rbind(tmpDF, data.frame(neighbors = i, avg.train.error = avg.train.error,
                                  avg.val.error = avg.val.error))

  cat("Average test error: ", avg.val.error, "With: ", i, " neighbors.\n")
}


## Average test error: 0.1600232 With: 2 neighbors.
## Average test error: 0.1233831 With: 10 neighbors.
## Average test error: 0.1201278 With: 20 neighbors.
## Average test error: 0.1209441 With: 30 neighbors.
## Average test error: 0.1193114 With: 40 neighbors.
## Average test error: 0.1209408 With: 50 neighbors.
## Average test error: 0.1229816 With: 60 neighbors.
## Average test error: 0.1270533 With: 70 neighbors.
## Average test error: 0.1233897 With: 80 neighbors.
## Average test error: 0.1225751 With: 90 neighbors.
## Average test error: 0.1213506 With: 100 neighbors.

tmpDF = tmpDF[-1,]

p1 <- qplot(neighbors, avg.train.error, data = tmpDF)
p2 <- qplot(neighbors, avg.val.error, data = tmpDF, colour= 'red')

grid.arrange(p1, p2, nrow = 1)

```



#Pick lowest val error for k in knn. #THIS IS DYNAMIC

```
Xtr = trn.cl[3:28]
Ytr = trn.cl$candidate
```

```
Xvl = tst.cl[3:28]
Yvl = tst.cl$candidate
```

```
predYtr = knn(train = Xtr, test = Xtr, cl = Ytr, k = 40)#might have to change k
```

```
predYvl = knn(train = Xtr, test = Xvl, cl = Ytr, k = 40)#might have to change k
```

```
records[2,]<-c(calc_error_rate(predYtr, Ytr),calc_error_rate(predYvl, Yvl))
records
```

```
##      train.error test.error
## tree  0.07247557 0.08469055
## knn   0.11482085 0.13355049
## lda           NA           NA
```

Classification: principal components

Instead of using the native attributes, we can use principal components in order to train our classification models. After this section, a comparison will be made between classification model performance between using native attributes and principal components.

```
pca.records = matrix(NA, nrow=3, ncol=2)
colnames(pca.records) = c("train.error", "test.error")
rownames(pca.records) = c("tree", "knn", "lda")
```


15. Compute principal components from the independent variables in training data. Then, determine the number of minimum number of PCs needed to capture 90% of the variance. Plot proportion of variance explained.

```
library(dplyr)
pca.records = matrix(NA, nrow=3, ncol=2)
colnames(pca.records) = c("train.error", "test.error")
rownames(pca.records) = c("tree", "knn", "lda")

trn.cl <- trn.cl %>% dplyr::select(-total)
PC.trn <- prcomp(trn.cl[2:27], scale=TRUE)

explained=0
values<-PC.trn$sdev^2
sum=sum(values)
y=NULL
for (i in 1:25){
  explained=explained + (values[i]/sum)
  y<-rbind(y,explained)}

plot(y, col=ifelse(y>=.90, "red", "black"), pch=ifelse(y>=.90, 19, 1), xlab='# of PCs', ylab='amount of variance explained')
```



16. Create a new training data by taking class labels and principal components. Call this variable `tr.pca`. Create the test data based on principal component loadings: i.e., transforming independent variables in test data to principal components space. Call this variable `test.pca`.

```
pca.records = matrix(NA, nrow=3, ncol=2)
colnames(pca.records) = c("train.error", "test.error")
rownames(pca.records) = c("tree", "knn", "lda")

tr.pca<-(PC.trn$x)
tr.pca<-as.data.frame(tr.pca)
tr.pca<-mutate(tr.pca, candidate = trn.cl$candidate)

test.pca <-(prcomp(tst.cl[,2:27], scale = TRUE)$x)
test.pca<-as.data.frame(test.pca)
test.pca<-mutate(test.pca, candidate=tst.cl$candidate)

length(test.pca)
```

```
## [1] 27
```

```
length(tr.pca)
```


```
## [1] 27
```

17. Decision tree: repeat training of decision tree models using principal components as independent variables. Record resulting errors.

```
tree.2<-tree(candidate~., data = tr.pca)
summary(tree.2)
```

```
##
## Classification tree:
## tree(formula = candidate ~ ., data = tr.pca)
## Variables actually used in tree construction:
## [1] "PC2" "PC1" "PC10" "PC3" "PC4" "PC15" "PC17"
## Number of terminal nodes: 12
## Residual mean deviance: 0.4685 = 1145 / 2444
## Misclassification error rate: 0.0908 = 223 / 2456

draw.tree(tree.2, nodeinfo=TRUE)
```

```

train.prediction.2<-predict(pruned.tree.2,tr.pca[-length(tr.pca)],type='class')
train.error.2<-calc_error_rate(train.prediction.2,tr.pca$candidate)

test.prediction.2<- predict(pruned.tree.2,test.pca[-length(test.pca)],type='class')
test.error.2<-calc_error_rate(test.prediction.2, test.pca$candidate)

#keeping track of errors
pca.records[1,]<-c(train.error.2,test.error.2)
pca.records

```

```

##      train.error test.error
## tree    0.110342  0.2149837
## knn           NA           NA
## lda           NA           NA

```

Interpretation & Discussion

19. This is an open question. Interpret and discuss any insights gained and possible explanations. Use any tools at your disposal to make your case: visualize errors on the map, discuss what does/doesn't seem reasonable based on your understanding of these methods, propose possible directions (collecting additional data, domain knowledge, etc)

To summarise we performed ten fold cross validation on the county census data set using the native attributes and the principal components as independent variables. We trained decision trees and KNN (K nearest neighbors) classifiers with the two variations of the datasets; the results shown:

```
records
```

```
##      train.error test.error
## tree  0.07247557 0.08469055
## knn   0.11482085 0.13355049
## lda           NA          NA
```

```
pca.records
```

```
##      train.error test.error
## tree    0.110342  0.2149837
## knn           NA          NA
## lda           NA          NA
```

Taking it further

20. Propose and tackle at least one interesting question. Be creative! Some possibilities are:

- Data preprocessing: we aggregated sub-county level data before performing classification. Would classification at the sub-county level before determining the winner perform better? What implicit assumptions are we making?
- Feature engineering: would a non-linear classification method perform better? Would you use native features or principal components?
- Additional classification methods: logistic regression, LDA, QDA, SVM, random forest, etc. (You may use methods beyond this course). How do these compare to KNN and tree method?
- Bootstrap: Perform bootstrap to generate plots similar to Figure 4.10/4.11. Discuss the results.

We use subcounty data to train a decision tree with native attributes and then compare it the tree trained for the aggregated data.

```
records2 = matrix(NA, nrow=1, ncol=2)
colnames(records2) = c("train.error", "test.error")
rownames(records2) = c("tree")

pca.records2 = matrix(NA, nrow=1, ncol=2)
colnames(pca.records2) = c("train.error", "test.error")
rownames(pca.records2) = c("tree")

#-----
tmpwinner = county_winner %>% ungroup %>%
  mutate(state = state.name[match(state, state.abb)]) %>%
  mutate_at(vars(state, county), tolower) %>%
  mutate(county = gsub(" county| columbia| city| parish", "", county))
tmpcensus = census.subct %>% mutate_at(vars(State, County), tolower)

election.cl = tmpwinner %>%
  left_join(tmpcensus, by = c("state"="State", "county"="County")) %>% na.omit

attr(election.cl, "location") = election.cl %>% dplyr::select(c(county, fips, state, votes, pct))

election.cl = election.cl %>% dplyr::select(-c(county, fips, state, votes, pct))

set.seed(69)
```

```
n = nrow(election.cl)
in.trn= sample.int(n, 0.8*n)
trn.cl = election.cl[ in.trn,]
tst.cl = election.cl[-in.trn,]
```